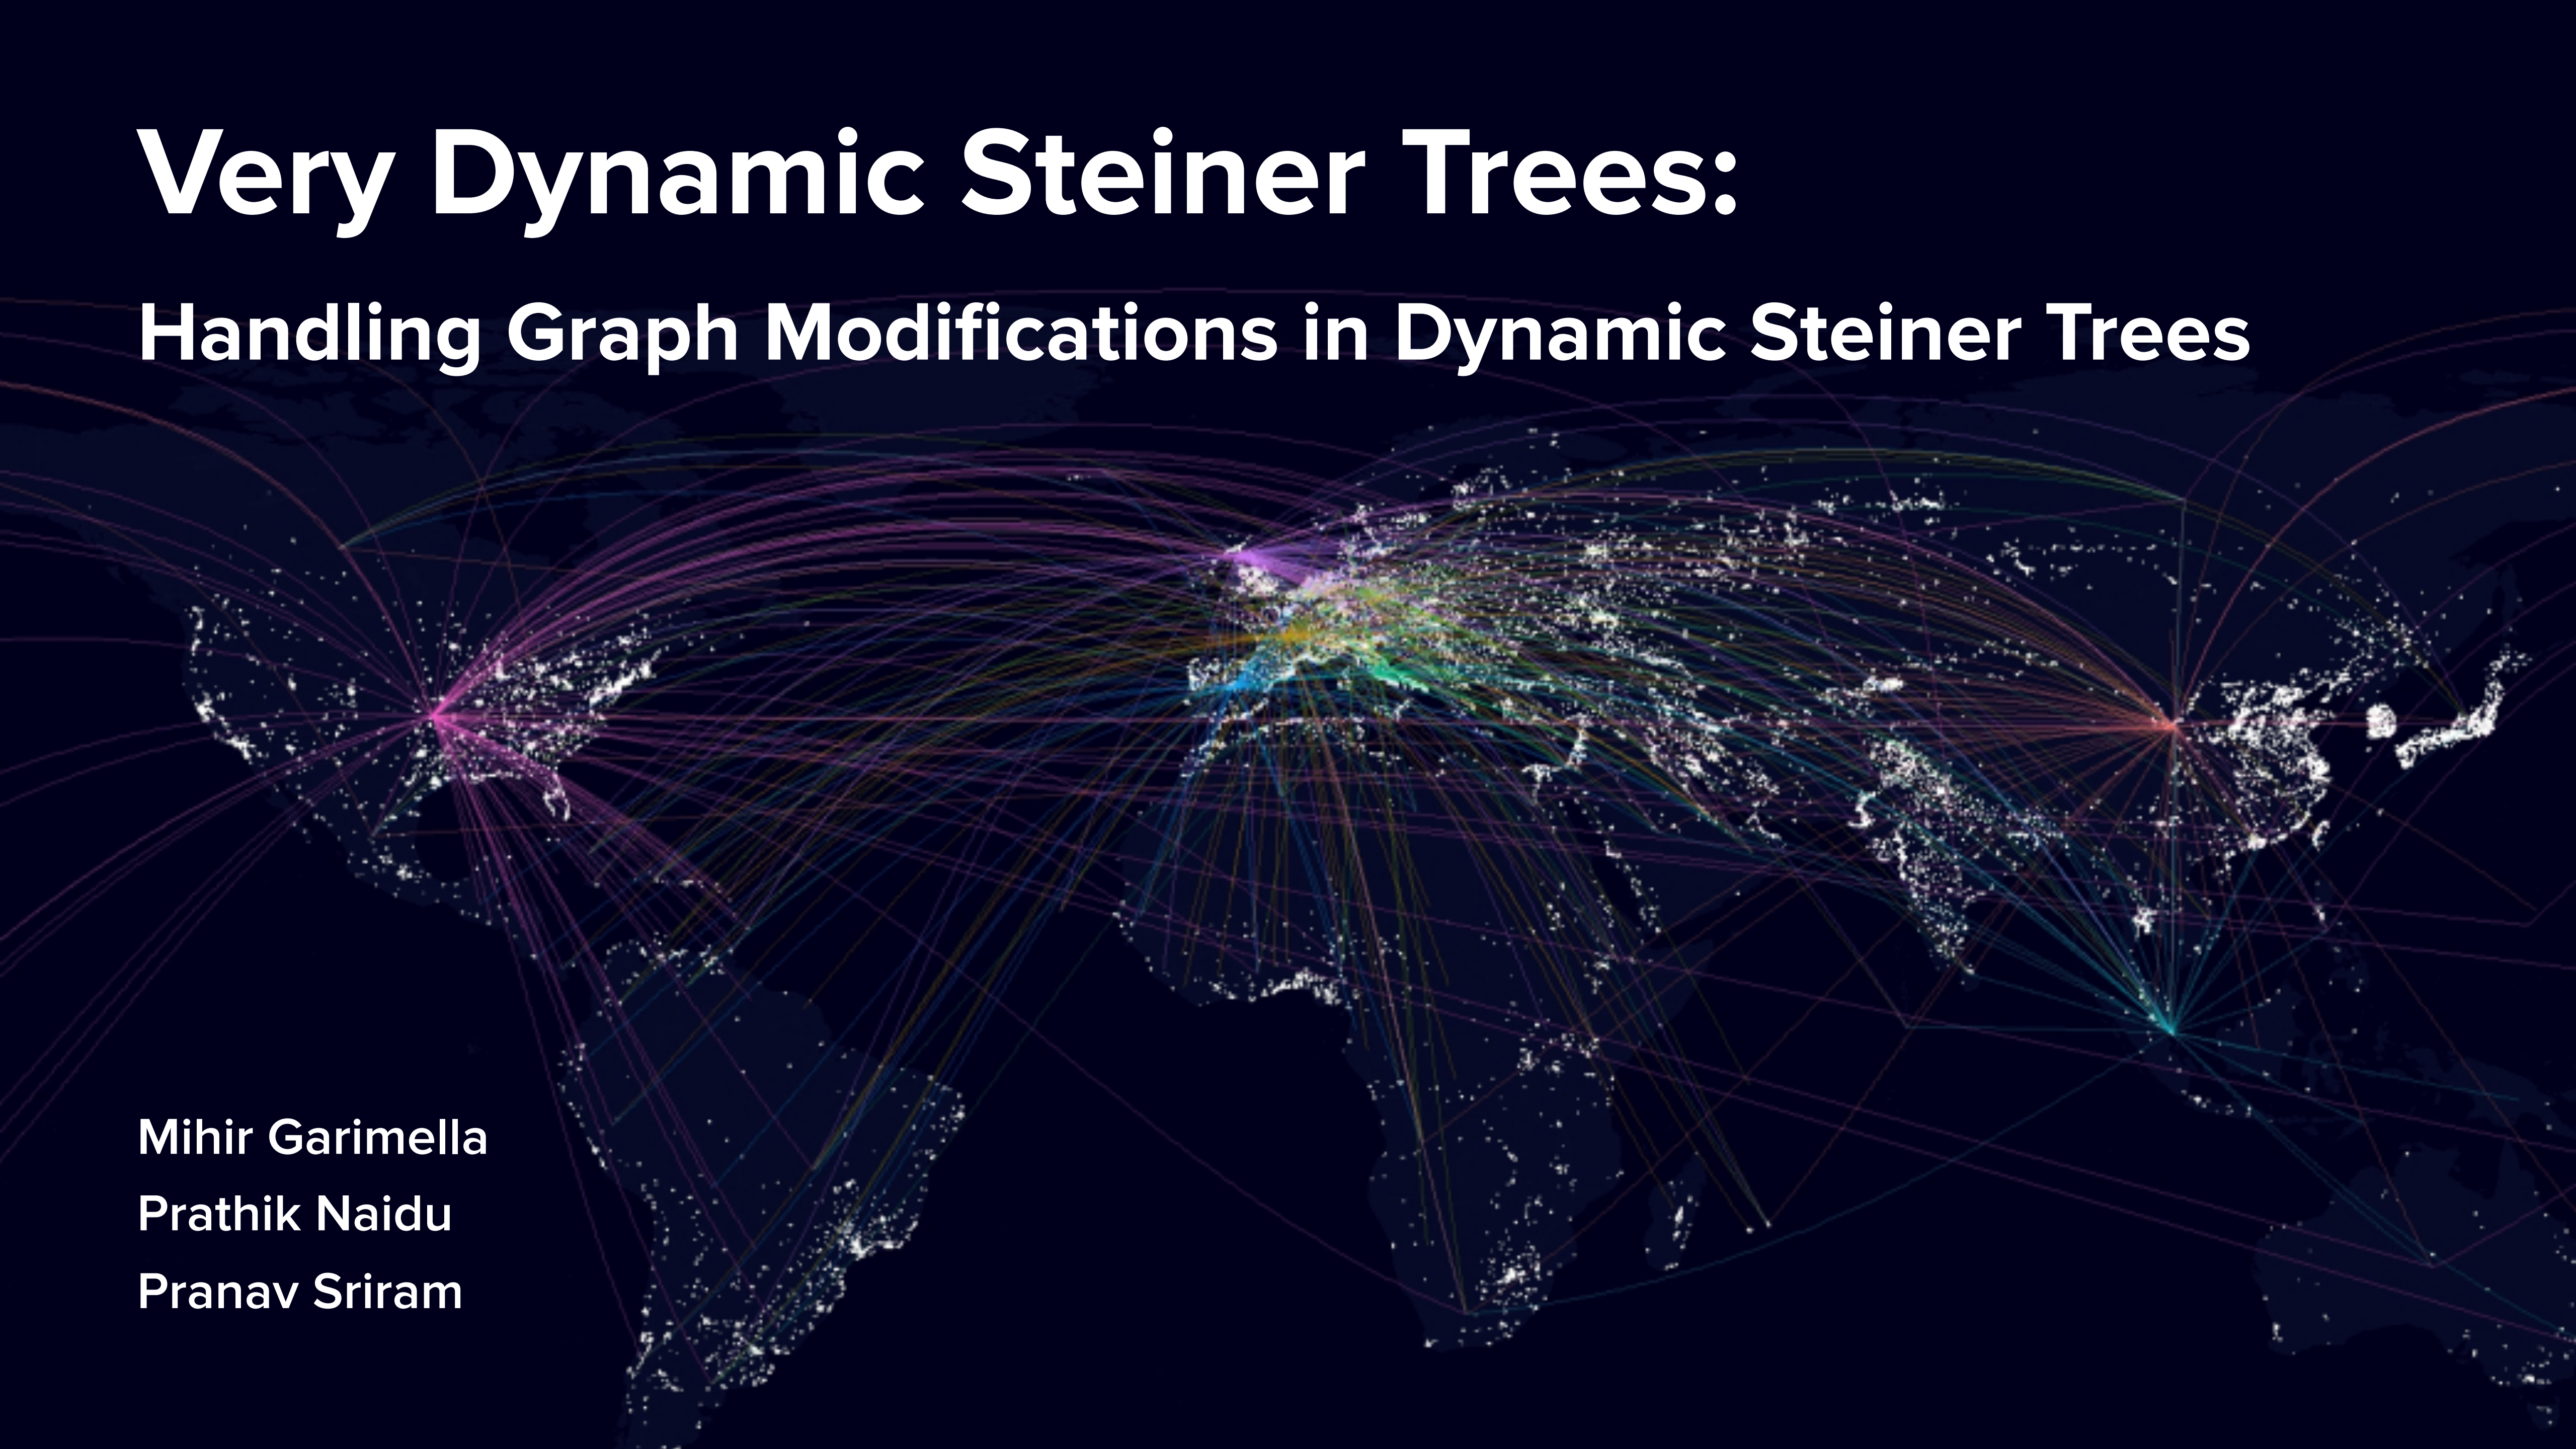


# Very Dynamic Steiner Trees:

## Handling Graph Modifications in Dynamic Steiner Trees

Mihir Garimella  
Prathik Naidu  
Pranav Sriram






# Introduction


**Q: What's the cheapest way to connect terminals?**

$t_i \in T$   
“Terminal” nodes  
we want to  
connect




# A: Just find the MST!

$t_i \in T$   
“Terminal” nodes  
we want to  
connect




# A: Just find the MST!


$t_i \in T$   
“Terminal” nodes  
we want to  
connect




**Q: What about now?**



**A: Well, we still want to connect all terminals...**




**...but it might be cheaper to use other nodes as well!**







**...but it might be cheaper to use other nodes as well!**




**...but it might be cheaper to use other nodes as well!**



**...but it might be cheaper to use other nodes as well!**




That's a *Steiner tree*!







**Q: Now, suppose we just start with a regular graph...**




**...and we're given terminals one-by-one?**



**...and we're given terminals one-by-one?**




**...and we're given terminals one-by-one?**







**...and we're given terminals one-by-one?**




**...and we're given terminals one-by-one?**



**...and we're given terminals one-by-one?**




**...and we're given terminals one-by-one?**







**...and we're given terminals one-by-one?**



**...and we're given terminals one-by-one?**




**...and we're given terminals one-by-one?**




# The Dynamic Steiner Tree (DST) Problem


Given a graph  $G = (V, E)$   
with positive edge costs...



...and terminals  $t_1, \dots, t_j \in T$   
arriving one-by-one,



maintain a subgraph  $T_j$  of  $G$   
that contains all  $t_i \in T$  and is  
close to the optimal Steiner  
tree after each new terminal.





# Intuitively...

Let's think about DSTs from the perspective of an adversary.



# Intuitively...

Let's think about DSTs from the perspective of an adversary.




# Intuitively...

Let's think about DSTs from the perspective of an adversary.




# Intuitively...

Let's think about DSTs from the perspective of an adversary.



# Intuitively...


Let's think about DSTs from the perspective of an adversary.






# Intuitively...

Let's think about DSTs from the perspective of an adversary.




# Intuitively...

Let's think about DSTs from the perspective of an adversary.



# Intuitively...

It's initially easy to place terminals far apart, but gets harder as more terminals arrive!



With many terminals, the cheapest path between a new terminal and any existing terminal probably won't be *that* expensive.

# Imase & Waxman (1991)

Proposed a simple yet effective **greedy algorithm**:

Add every new terminal by connecting it via the cheapest path to any of the existing terminals in the graph.

$2 \log k$  competitive, where  $k$  is the number of terminals

**Key limitation:** only really handles terminal addition

# Our Work



# VDSTGraph

Graph data structure that supports **five operations**:

**Key idea:** do some extra work here to make the fifth operation fast

1 Add node

2 Add edge (or decrease weight)

3 Add terminal ] (The only one of these that Imase & Waxman addressed!)

4 Remove terminal

5 Query Steiner tree

# The *Metric* Steiner Tree Problem

“Metric” =

- 1 Graph is complete
- 2 Edge weights satisfy Triangle Inequality

**Theorem:** Any  $\alpha$ -competitive algorithm for the **metric Steiner tree problem** is also  $\alpha$ -competitive for the general Steiner tree problem.

**How?** Find the Steiner tree on the **metric closure**, then expand edges into full paths and break cycles.

# How VDSTGraph Works

**Intuition:** Easy to handle graph updates in metric problem, but not so much in the general case.

So, we efficiently maintain the **metric closure** and **its Steiner tree** and lazily generalize the results to our original graph at query time.

To keep our data structure  $2 \log k$  competitive, we maintain the “**favorite parent invariant**” during graph updates.

Convention:  $n = \#$  of nodes,  $m = \#$  of edges,  $k = \#$  of terminals

# VDSTGraph.add\_node

Easy — node is initially disconnected from others and not a terminal.

- 1 Add node to graph and metric closure
- 2 Create direct path to all other nodes in metric closure with cost  $\infty$

Time complexity:  $O(n)$

# VDSTGraph.decrease\_edge\_weight

Harder — might introduce many shortest paths around the graph.

- 1 Add this edge to metric closure
- 2 Using DP, solve *All Pairs Shortest Path* problem to “propagate” effect in metric closure
- 3 Update metric Steiner tree each time whenever we find a shorter path between two terminals

Time complexity:  $O(n^2)$

# VDSTGraph.add\_terminal

Easy — just do what Imase & Waxman did!

- 1 Add node to metric Steiner tree
- 2 Using metric closure, find and connect to closest terminal

Time complexity:  $O(k)$



# VDSTGraph.remove\_terminal

Harder — might cause other terminals to be orphaned.

- 1 Remove node from metric Steiner tree
- 2 For each child, reattach with add\_terminal (respect timestamps!)

Time complexity:  $O(k^2)$

# VDSTGraph.get\_steiner\_tree

Easy — just rebuild general Steiner tree from metric one.

- 1 Expand paths in metric Steiner tree, adding edges to priority queue.
- 2 Run Kruskal's algorithm to find lowest-cost way to break cycles.

Time complexity:  $O(m \log n)$

Cache queries to avoid recomputation if nothing changes.

# Results

# Evaluation Pipeline

## Baseline: NetworkX Python package

```
steiner_tree(G, terminal_nodes, weight='weight') \[source\]
```

Return an approximation to the minimum Steiner tree of a graph.

$(2 - (2/k))$  competitive

- 1 Build initial graph with 300 initial nodes and edge density  $\rho$
- 2 Test both methods on randomized workflow of 500 operations

**Add terminal**

then...

**Query Steiner tree**

**Add terminal**

**Add node + edges**

**Remove terminal**


**Decrease edge weight**

(first 40 operations)

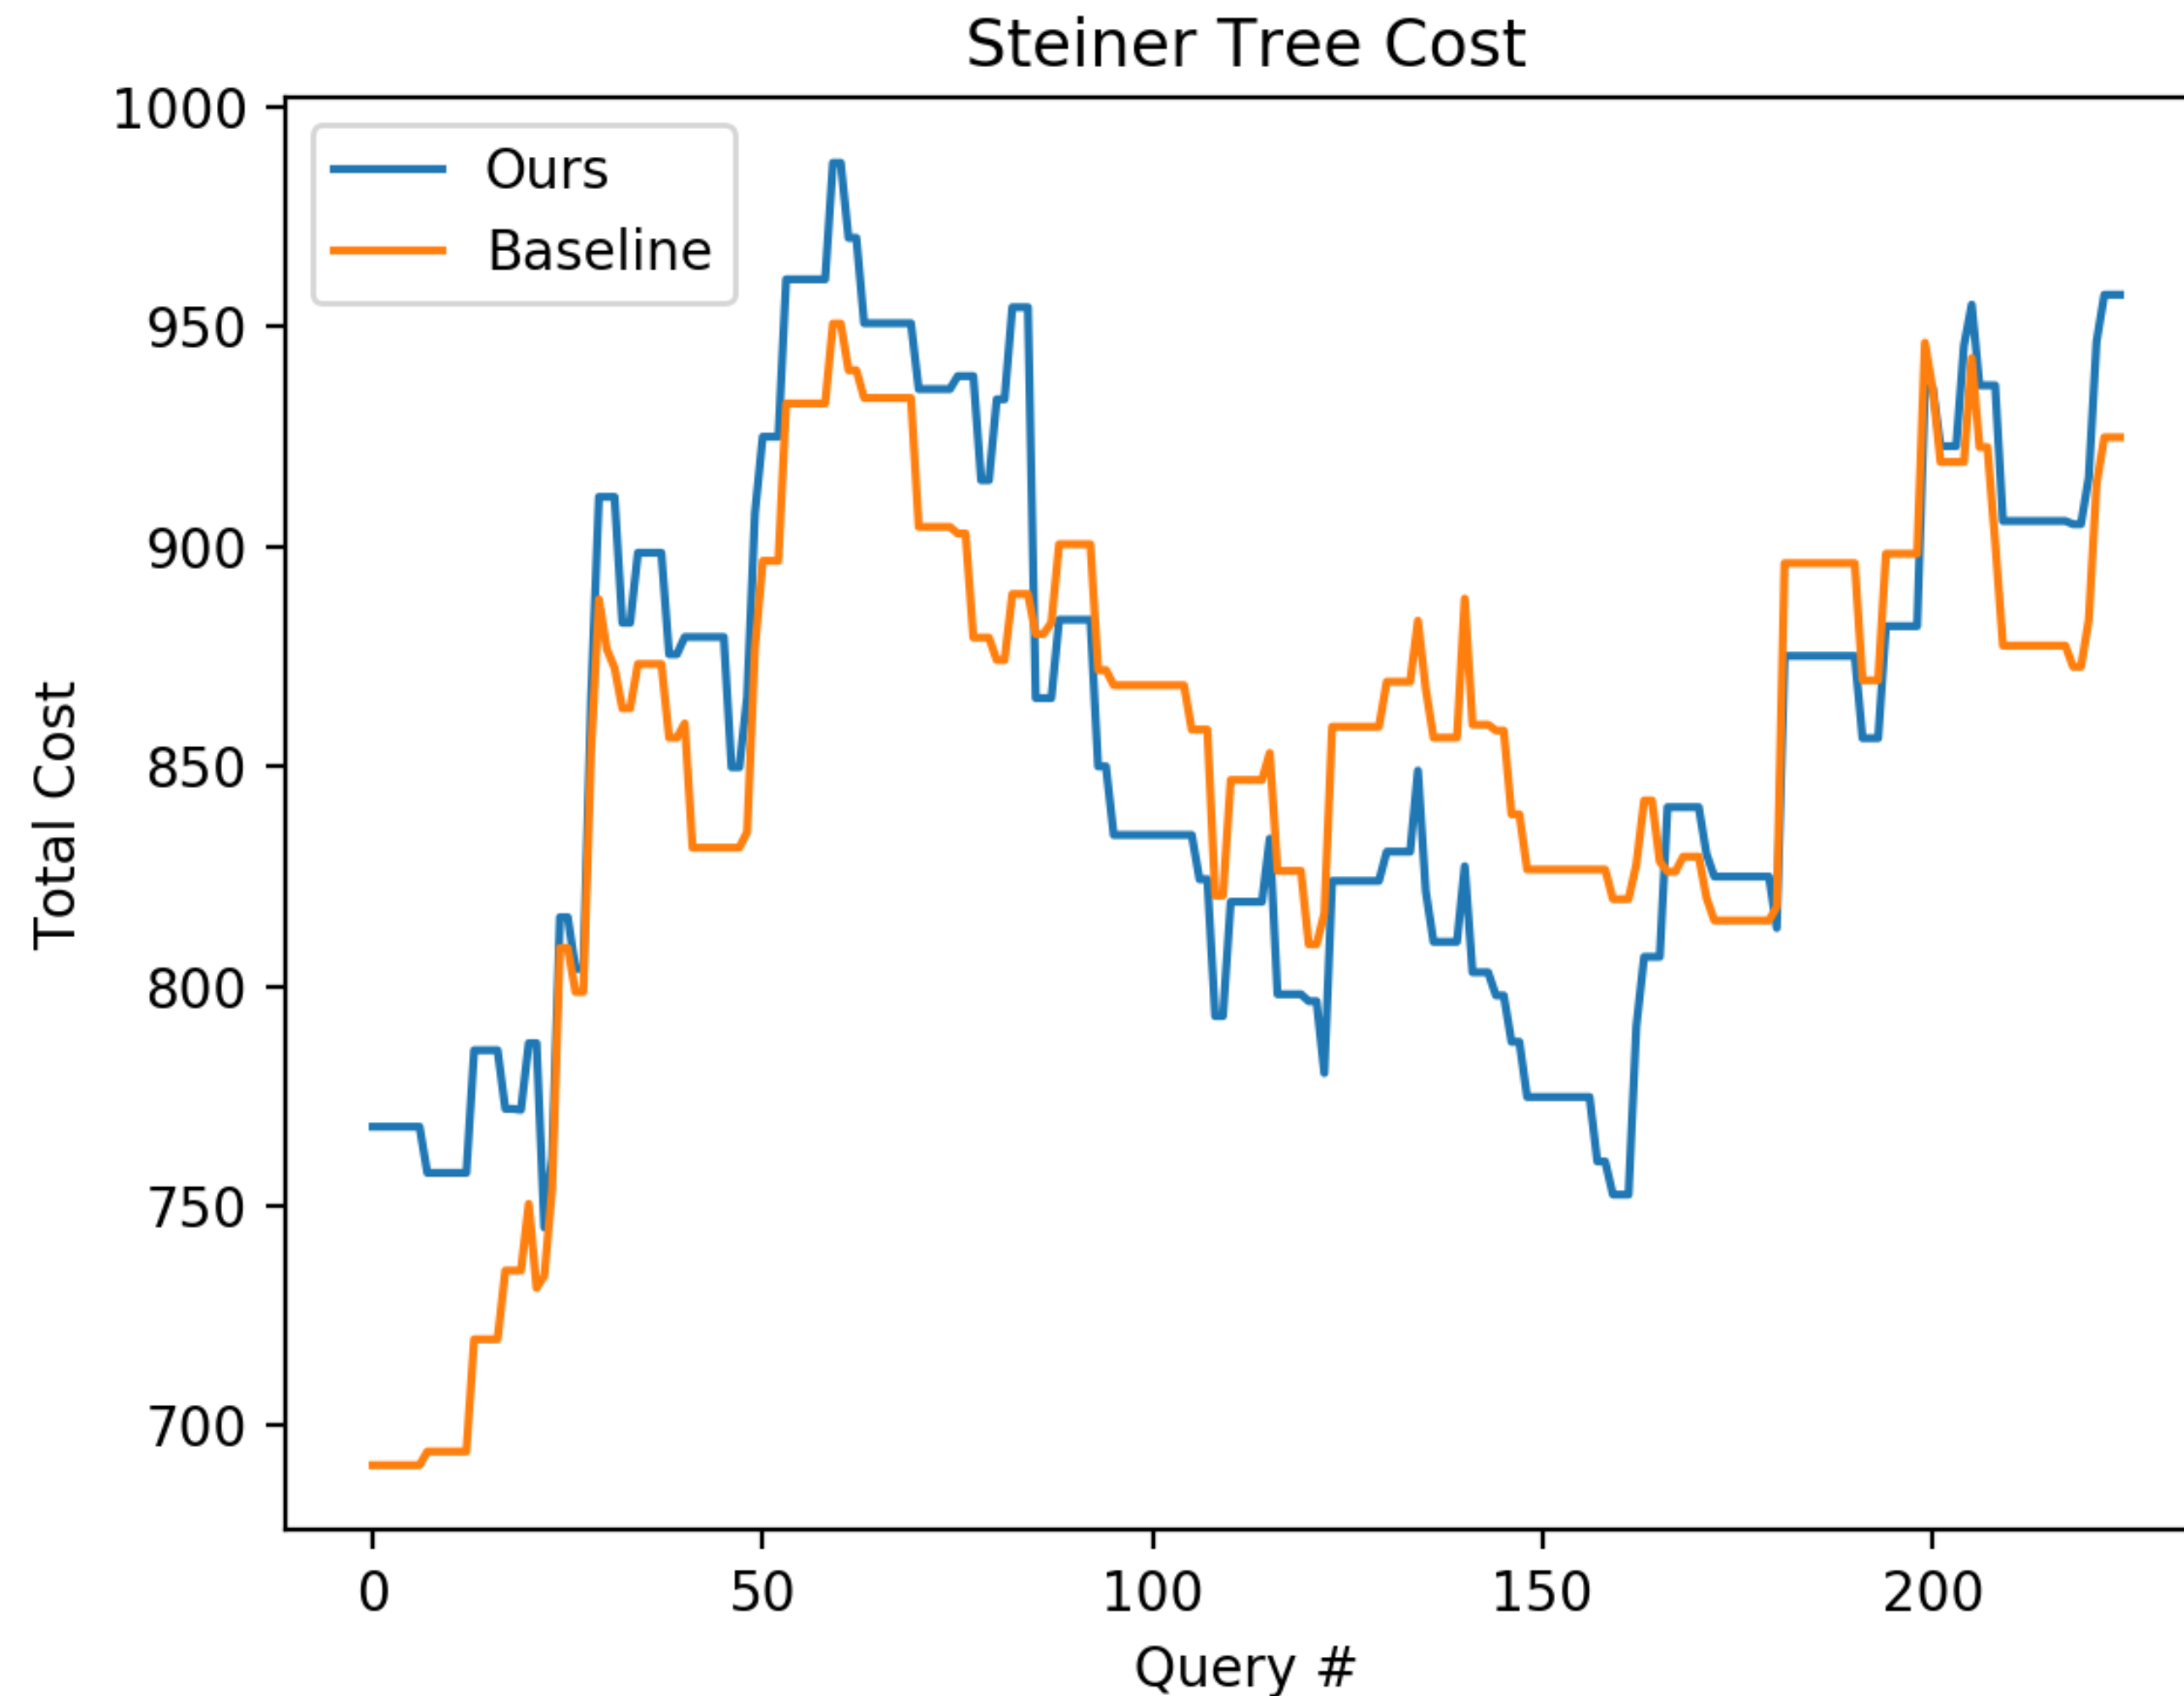
(with probability  $p$ )

(with equal probability if we don't query)

# Q: Does it work?



# A: Yes!





No noticeable reduction in quality compared to baseline (a more accurate approximation)!


On average, returned Steiner trees with **1.049** times cost of those returned by baseline.




# Q: Is it fast?



# A: Extremely!




Faster as workflow includes more queries!



Scales extremely well to dense graphs.

# Who cares?

Generalizing graph algorithms to the dynamic case is non-trivial



**Theory**

DSTs used in dynamic routing (Uber), large-scale distribution (Amazon), etc.



**Practice**

# Future Work

- 1 Generalize metric Steiner tree without *de novo* recomputation at query time (path expansion, Kruskal's)
- 2 Handle node and edge deletion
- 3 Handle batch insertions/deletions
- 4 Evaluate on real-world datasets (Oregon Networking or U.S. Road Networks)

# Conclusion

**To our knowledge, first to address graph modifications in the DST problem.**

Algorithm handles dynamic modifications to terminal nodes, graph nodes, and edges.

No loss in accuracy, but significantly faster than existing algorithms.

Scales extremely well to dense graphs/lots of queries by **reusing more and more work**.